

Appendix - San Diego Street Conditions Classification

A Cloud Computing Project by Leonid Shpaner, Jose Luis Estrada, and Kiran Singh

```
[1]: import boto3, re, sys, math, json, os, sagemaker, urllib.request
import io
import sagemaker
from sagemaker import get_execution_role
from IPython.display import Image
from IPython.display import display
from time import gmtime, strftime
from sagemaker.predictor import csv_serializer
from pyathena import connect
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from prettytable import PrettyTable
from imblearn.over_sampling import SMOTE, ADASYN
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split, \
RepeatedStratifiedKFold, RandomizedSearchCV
from sklearn.metrics import roc_curve, auc, mean_squared_error, \
precision_score, recall_score, f1_score, accuracy_score, \
confusion_matrix, plot_confusion_matrix, classification_report
from sagemaker.tuner import HyperparameterTuner
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from scipy.stats import loguniform
import warnings
warnings.filterwarnings('ignore')
```

Data Wrangling

```
[2]: # create athena database
sess = sagemaker.Session()
bucket = sess.default_bucket()
role = sagemaker.get_execution_role()
region = boto3.Session().region_name
# s3 = boto3.Session().client(service_name="s3", region_name=region)

# ec2 = boto3.Session().client(service_name="ec2", region_name=region)
# sm = boto3.Session().client(service_name="sagemaker", region_name=region)
```

```
[3]: ingest_create_athena_db_passed = False
```

```
[4]: # set a database name
database_name = "watersd"
```

```
[5]: # Set S3 staging directory -- this is a temporary directory used for Athena queries
s3_staging_dir = "s3://{0}/athena/staging".format(bucket)
```

```
[6]: conn = connect(region_name=region, s3_staging_dir=s3_staging_dir)
```

```
[7]: statement = "CREATE DATABASE IF NOT EXISTS {}".format(database_name)
print(statement)
pd.read_sql(statement, conn)
```

```
CREATE DATABASE IF NOT EXISTS watersd
```

```
[7]: Empty DataFrame
Columns: []
Index: []
```

```
[8]: water_dir = 's3://waterteam1/raw_files'
```

```
[9]: # SQL statement to execute the analyte tests drinking water table

table_name = 'oci_2015_datasd'
pd.read_sql(f'DROP TABLE IF EXISTS {database_name}.{table_name}', conn)

create_table = f"""
CREATE EXTERNAL TABLE IF NOT EXISTS {database_name}.{table_name}(
    seg_id string,
    oci float,
    street string,
    street_from string,
    street_to string,
    seg_length_ft float,
    seg_width_ft float,
    func_class string,
    pvm_class string,
    area_sq_ft float,
    oci_desc string,
    oci_wt float
)

ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '{water_dir}/{table_name}'
TBLPROPERTIES ('skip.header.line.count'='1')
"""

pd.read_sql(create_table, conn)

pd.read_sql(f'SELECT * FROM {database_name}.{table_name} LIMIT 5', conn)
```

```
[9]:
```

	seg_id	oci	street	street_from	street_to	seg_length_ft	seg_width_ft	\
0	SA-000003	65.14	ALLEY			772.7258	30.0	
1	SA-000004	67.45	ALLEY			196.0025	30.0	
2	SA-000005	70.88	ALLEY			395.0049	30.0	
3	SA-000006	84.00	ALLEY			192.0025	30.0	
4	SA-000008	79.24	ALLEY			251.7540	30.0	

	func_class		pvm_class	area_sq_ft	oci_desc	oci_wt
0	Alley	PCC	Jointed Concrete	23181.773	Fair	1510060.80
1	Alley	PCC	Jointed Concrete	5880.075	Fair	396611.06
2	Alley	PCC	Jointed Concrete	11850.147	Good	839938.44
3	Alley	PCC	Jointed Concrete	5760.075	Good	483846.30
4	Alley	PCC	Jointed Concrete	7552.620	Good	598469.60

[10]: *# SQL statement to execute the analyte tests drinking water table*

```
table_name2 = 'sd_paving_dataasd'
pd.read_sql(f'DROP TABLE IF EXISTS {database_name}.{table_name2}', conn)

create_table = f"""
CREATE EXTERNAL TABLE IF NOT EXISTS {database_name}.{table_name2}(
    pve_id int,
    seg_id string,
    project_id string,
    title string,
    project_manager string,
    project_manager_phone string,
    status string,
    type string,
    resident_engineer string,
    address_street string,
    street_from string,
    street_to string,
    seg_cd int,
    length int,
    width int,
    date_moratorium date,
    date_start date,
    date_end date,
    paving_miles float
)

ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '{water_dir}/{table_name2}'
TBLPROPERTIES ('skip.header.line.count'='1')
"""

pd.read_sql(create_table, conn)

pd.read_sql(f'SELECT * FROM {database_name}.{table_name2} LIMIT 5', conn)
```

[10]:

	pve_id	seg_id	project_id	title	\
0	1073577074	SA-000319	UTLY	Public Works CIP	
1	1792486183	SA-000345	UTLY	Public Works CIP	
2	1173780646	SA-000375	UTLY	Public Works CIP	

3	1276790298	SA-000378	UTLY	Public Works	CIP
4	27170959	SA-001081	UTLY	Public Works	CIP

	project_manager	project_manager_phone	status	\
0	Engineering@sandiego.gov	858-627-3200	post construction	
1	Engineering@sandiego.gov	858-627-3200	post construction	
2	Engineering@sandiego.gov	858-627-3200	post construction	
3	Engineering@sandiego.gov	858-627-3200	post construction	
4	Engineering@sandiego.gov	858-627-3200	post construction	

	type	resident_engineer	address_street	street_from	street_to	seg_cd	\
0	Overlay	ECP	ALLEY			2	
1	Slurry	ECP	ALLEY			2	
2	Slurry	ECP	ALLEY			2	
3	Slurry	ECP	ALLEY			2	
4	Concrete	ECP	ALLEY			9	

	length	width	date_moratorium	date_start	date_end	paving_miles
0	0	NaN	2019-02-02	2019-02-02	2019-02-02	0.000000
1	938	30.0	2019-01-30	2019-01-30	2019-01-30	0.177652
2	674	30.0	2018-08-01	2018-08-01	2018-08-01	0.127652
3	658	30.0	2018-08-01	2018-08-01	2018-08-01	0.124621
4	680	30.0	None	2020-08-13	2020-08-13	0.128788

```
[11]: # SQL statement to execute the analyte tests drinking water table

table_name3 = 'traffic_counts_datasd'
pd.read_sql(f'DROP TABLE IF EXISTS {database_name}.{table_name3}', conn)

create_table = f"""
CREATE EXTERNAL TABLE IF NOT EXISTS {database_name}.{table_name3}(
    id string,
    street_name string,
    limits string,
    northbound_count int,
    southbound_count int,
    eastbound_count int,
    westbound_count int,
    total_count int,
    file_no string,
    date_count date
)

ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '{water_dir}/{table_name3}'
TBLPROPERTIES ('skip.header.line.count'='1')
"""

pd.read_sql(create_table, conn)
```

```
pd.read_sql(f'SELECT * FROM {database_name}.{table_name3} LIMIT 5', conn)
```

```
[11]:
```

	id	street_name	limits	northbound_count	\
0	01AV018207	01 AV	A ST - ASH ST	18010	
1	01AV015210	01 AV	A ST - ASH ST	20060	
2	01AV018213	01 AV	A ST - ASH ST	19597	
3	01AV007721	01 AV	A ST - ASH ST	10640	
4	01AV088812	01 AV	ASH ST - BEECH ST	2298	

	southbound_count	eastbound_count	westbound_count	total_count	file_no	\
0	None	None	None	18010	0182-07	
1	None	None	None	20060	0152-10	
2	None	None	None	19597	0182-13	
3	None	None	None	10640	0077-21	
4	None	None	None	2298	0888-12	

	date_count
0	2007-03-13
1	2010-03-18
2	2013-03-12
3	2021-03-10
4	2012-12-11

```
[12]: statement = "SHOW DATABASES"
df_show = pd.read_sql(statement, conn)
df_show.head(5)
```

```
[12]:
```

	database_name
0	default
1	dsoaws
2	watersd

```
[13]: if database_name in df_show.values:
       ingest_create_athena_db_passed = True
```

```
[14]: %store ingest_create_athena_db_passed
```

Stored 'ingest_create_athena_db_passed' (bool)

```
[15]: pd.read_sql(f'SELECT * FROM {database_name}.{table_name} t1 INNER JOIN \
                {database_name}.{table_name2} t2 ON t1.seg_id \
                = t2.seg_id LIMIT 5', conn)
```

```
[15]:
```

	seg_id	oci	street	street_from	street_to	seg_length_ft	\
0	SA-000345	34.14	ALLEY			937.9261	
1	SA-000375	97.25	ALLEY			673.3209	
2	SA-000378	62.67	ALLEY			657.2000	
3	SA-001081	68.86	ALLEY			679.1060	
4	SA-001083	28.67	ALLEY			660.0917	

	seg_width_ft	func_class	pvm_class	area_sq_ft	...	\
--	--------------	------------	-----------	------------	-----	---

0	30.0	Alley	AC Improved	28137.783	...
1	30.0	Alley	PCC Jointed Concrete	20199.627	...
2	30.0	Alley	PCC Jointed Concrete	19716.000	...
3	30.0	Alley	PCC Jointed Concrete	20373.180	...
4	30.0	Alley	PCC Jointed Concrete	19802.752	...

	address_street	street_from	street_to	seg_cd	length	width	date_moratorium	\
0	ALLEY			2	938	30	2019-01-30	
1	ALLEY			2	674	30	2018-08-01	
2	ALLEY			2	658	30	2018-08-01	
3	ALLEY			9	680	30	None	
4	ALLEY			9	661	30	None	

	date_start	date_end	paving_miles
0	2019-01-30	2019-01-30	0.177652
1	2018-08-01	2018-08-01	0.127652
2	2018-08-01	2018-08-01	0.124621
3	2020-08-13	2020-08-13	0.128788
4	2020-07-31	2020-07-31	0.125189

[5 rows x 31 columns]

```
[16]: df = pd.read_sql(f'SELECT * FROM (SELECT * FROM {database_name}.{table_name} \
                        t1 INNER JOIN {database_name}.{table_name2} t2 \
                        ON t1.seg_id = t2.seg_id) m1 LEFT JOIN (SELECT street_name, \
                        SUM(total_count) \
                        FROM \
                        {database_name}.{table_name3} \
                        GROUP BY \
                        street_name) t3 \
                        ON m1.address_street = t3.street_name', conn)
```

```
[17]: df.head(5)
```

[17]:	seg_id	oci	street	street_from	street_to	seg_length_ft	seg_width_ft	\
0	SA-000345	34.14	ALLEY			937.9261	30.0	
1	SA-000375	97.25	ALLEY			673.3209	30.0	
2	SA-000378	62.67	ALLEY			657.2000	30.0	
3	SA-001081	68.86	ALLEY			679.1060	30.0	
4	SA-001083	28.67	ALLEY			660.0917	30.0	

	func_class	pvm_class	area_sq_ft	...	street_to	seg_cd	length	\
0	Alley	AC Improved	28137.783	...		2	938	
1	Alley	PCC Jointed Concrete	20199.627	...		2	674	
2	Alley	PCC Jointed Concrete	19716.000	...		2	658	
3	Alley	PCC Jointed Concrete	20373.180	...		9	680	
4	Alley	PCC Jointed Concrete	19802.752	...		9	661	

	width	date_moratorium	date_start	date_end	paving_miles	street_name	\
--	-------	-----------------	------------	----------	--------------	-------------	---

0	30	2019-01-30	2019-01-30	2019-01-30	0.177652	None
1	30	2018-08-01	2018-08-01	2018-08-01	0.127652	None
2	30	2018-08-01	2018-08-01	2018-08-01	0.124621	None
3	30	None	2020-08-13	2020-08-13	0.128788	None
4	30	None	2020-07-31	2020-07-31	0.125189	None

	total_count
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 33 columns]

```
[18]: # remove duplicated columns
df = df.loc[:, ~df.columns.duplicated()]
```

```
[19]: # create flat .csv file from originally
# merged dataframe
# df.to_csv('original_merge.csv')
```

Exploratory Data Analysis (EDA)

```
[20]: # get number of rows and columns
print('Number of Rows:', df.shape[0])
print('Number of Columns:', df.shape[1], '\n')

# inspect datatypes and nulls
data_types = df.dtypes
data_types = pd.DataFrame(data_types)
data_types = data_types.assign(Null_Values =
                               df.isnull().sum())
data_types.reset_index(inplace = True)
data_types.rename(columns={0: 'Data Type',
                           'index': 'Column/Variable',
                           'Null_Values': "# of Nulls"})
```

Number of Rows: 23005
Number of Columns: 30

[20]:	Column/Variable	Data Type	# of Nulls
0	seg_id	object	0
1	oci	float64	0
2	street	object	0
3	street_from	object	0
4	street_to	object	0
5	seg_length_ft	float64	0
6	seg_width_ft	float64	0
7	func_class	object	0
8	pvm_class	object	0

9	area_sq_ft	float64	0
10	oci_desc	object	0
11	oci_wt	float64	0
12	pve_id	int64	0
13	project_id	object	0
14	title	object	0
15	project_manager	object	0
16	project_manager_phone	object	0
17	status	object	0
18	type	object	0
19	resident_engineer	object	0
20	address_street	object	0
21	seg_cd	int64	0
22	length	int64	0
23	width	int64	0
24	date_moratorium	object	4426
25	date_start	object	1
26	date_end	object	7
27	paving_miles	float64	0
28	street_name	object	16874
29	total_count	float64	16874

Bias Exploration

To explore potential areas of bias, we will endeavor to trace class imbalance on the target feature of “oci_desc.”

```
[21]: oci_desc_fair = df['oci_desc'].value_counts()['Fair']
oci_desc_good = df['oci_desc'].value_counts()['Good']
oci_desc_poor = df['oci_desc'].value_counts()['Poor']
oci_desc_total = oci_desc_fair + oci_desc_good + oci_desc_poor

table1 = PrettyTable() # build a table
table1.field_names = ['Fair Condition', 'Good Condition',
                     'Poor Condition', 'Total']
table1.add_row([oci_desc_fair, oci_desc_good, oci_desc_poor,
                oci_desc_total])
table1
```

```
[21]: +-----+-----+-----+-----+
| Fair Condition | Good Condition | Poor Condition | Total |
+-----+-----+-----+-----+
|      6105     |      15758     |       1142     | 23005 |
+-----+-----+-----+-----+
```

```
[22]: perc_good = oci_desc_good / (oci_desc_total)
perc_fair = oci_desc_fair / (oci_desc_total)
perc_poor = oci_desc_poor / (oci_desc_total)
print(round(perc_good, 2)*100, '% of streets '
        'are in good condition ')
print(round(perc_fair, 2)*100, '% of streets '
        'are in fair condition ')
```



```
print(round(perc_poor, 2)*100, '% of streets '
      'are in poor condition ')
```

68.0 % of streets are in good condition

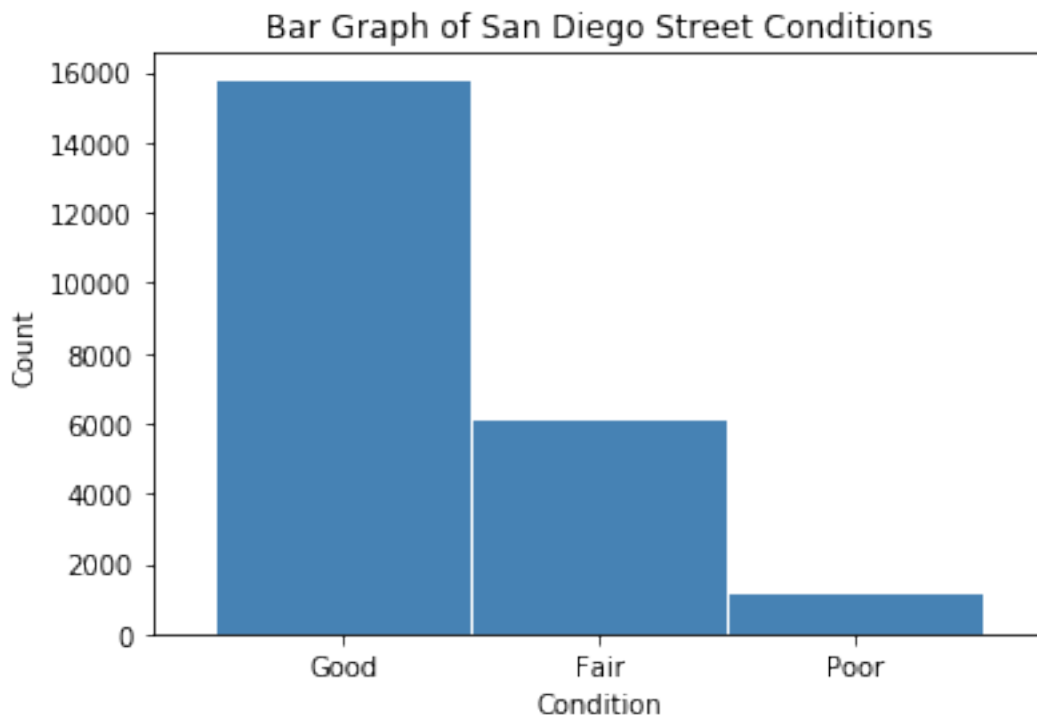
27.0 % of streets are in fair condition

5.0 % of streets are in poor condition

Considerably more than half of the streets are in good condition. A little less than a third are in fair condition. Only 5% are in poor condition.

```
[23]: # accidents injury bar graph
conditions = df['oci_desc'].value_counts()
fig = plt.figure()
conditions.plot.bar(x='lab', y='val', rot=0, width=0.99,
                  color="steelblue")
plt.title('Bar Graph of San Diego Street Conditions')
plt.xlabel('Condition')
plt.ylabel('Count')
plt.show()

conditions
```



```
[23]: Good      15758
      Fair       6105
      Poor       1142
      Name: oci_desc, dtype: int64
```

Whereas a method can be used to classify street conditions into multiple classes, it is easier to re-classify streets in “fair” and “good” condition into one category in comparison with the poor class. This, in turn,

becomes a binary classification problem. Thus, there are now 21,863 streets in good condition and 1,142 in poor condition (only 5% of all streets). This presents a definitive example of class imbalance.

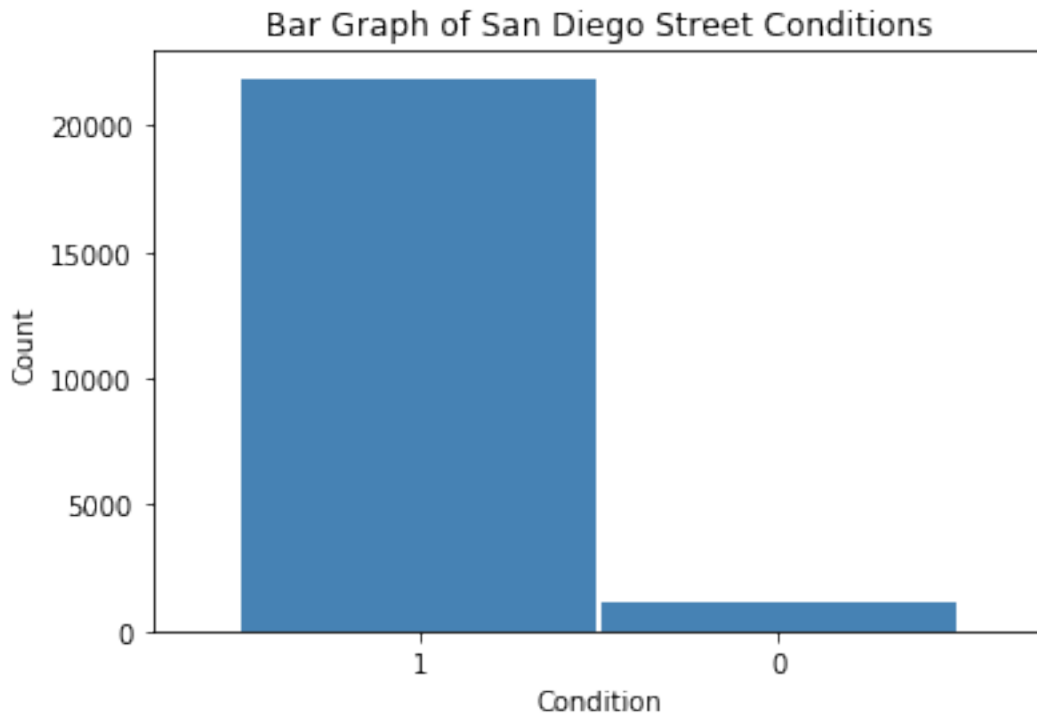
```
[24]: df['oci_cat'] = df['oci_desc'].map({'Good':1, 'Fair':1,
                                         'Poor':0})

cond = df['oci_cat'].value_counts()
cond
```

```
[24]: 1    21863
      0     1142
      Name: oci_cat, dtype: int64
```

```
[25]: # oci ratings bar graph
fig = plt.figure()
cond.plot.bar(x='lab', y='val', rot=0, width=0.99,
              color="steelblue")
plt.title('Bar Graph of San Diego Street Conditions')
plt.xlabel('Condition')
plt.ylabel('Count')
plt.show()

cond
```



```
[25]: 1    21863
      0     1142
      Name: oci_cat, dtype: int64
```

```
[26]: # cast oci info into range of values
labels = [ "{0} - {1}".format(i, i + 5) for i in range(0, 100, 10) ]
df['OCI Range'] = pd.cut(df.oci, range(0, 105, 10),
                        right=False,
                        labels=labels).astype(object)

# inspect the new dataframe with this info
df[['oci', 'OCI Range']]
```

```
[26]:
```

	oci	OCI Range
0	34.14	30 - 35
1	97.25	90 - 95
2	62.67	60 - 65
3	68.86	60 - 65
4	28.67	20 - 25
...
23000	93.40	90 - 95
23001	91.01	90 - 95
23002	97.26	90 - 95
23003	95.00	90 - 95
23004	80.83	80 - 85

[23005 rows x 2 columns]

```
[27]: print("\033[1m" + 'Street Conditions by Condition Index Range' + "\033[1m")
def oci_cond():
    oci_desc_good = df.loc[df.oci_desc == 'Good'].groupby(
        ['OCI Range'])[['oci_desc']].count()
    oci_desc_good.rename(columns = {'oci_desc': 'Good'}, inplace=True)
    oci_desc_fair = df.loc[df.oci_desc == 'Fair'].groupby(
        ['OCI Range'])[['oci_desc']].count()
    oci_desc_fair.rename(columns = {'oci_desc': 'Fair'}, inplace=True)
    oci_desc_poor = df.loc[df.oci_desc == 'Poor'].groupby(
        ['OCI Range'])[['oci_desc']].count()
    oci_desc_poor.rename(columns = {'oci_desc': 'Poor'}, inplace=True)
    oci_desc_comb = pd.concat([oci_desc_good, oci_desc_fair, oci_desc_poor],
        axis = 1)
    # sum row totals
    oci_desc_comb.loc['Total'] = oci_desc_comb.sum(numeric_only=True, axis=0)
    # sum column totals
    oci_desc_comb.loc[:, 'Total'] = oci_desc_comb.sum(numeric_only=True, axis=1)
    oci_desc_comb.fillna(0, inplace = True)
    return oci_desc_comb.style.format("{:,.0f}")

oci_cond = oci_cond().data # retrieve dataframe
oci_cond
```

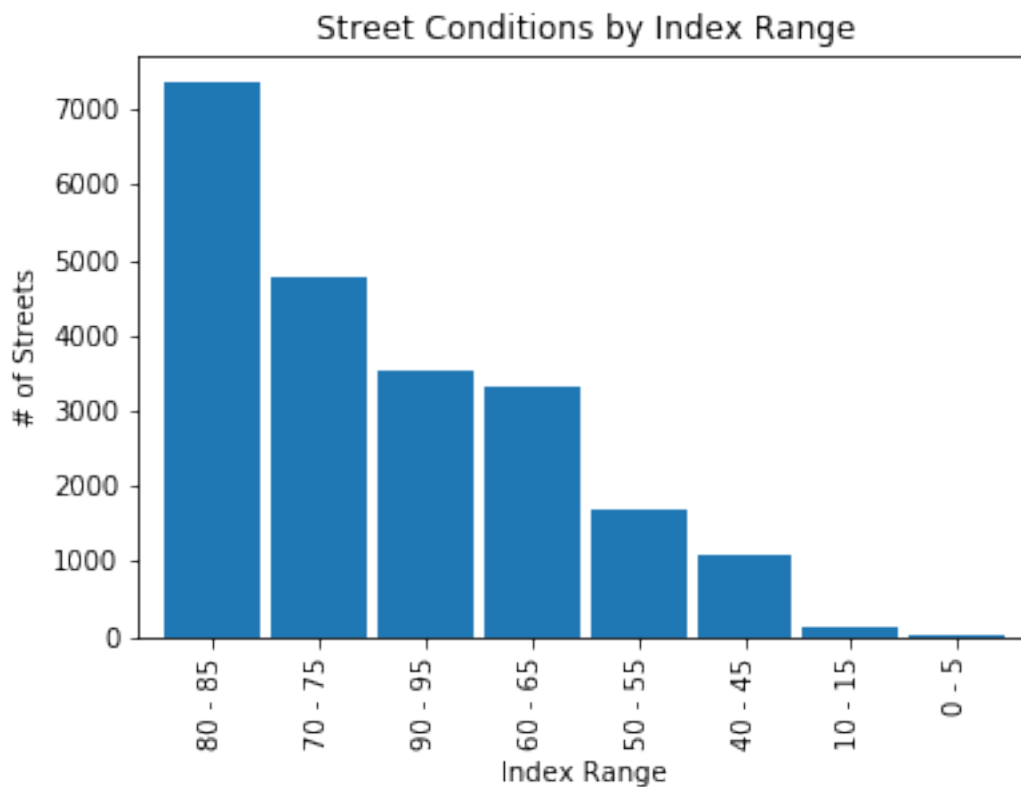
Street Conditions by Condition Index Range

```
[27]:
```

	Good	Fair	Poor	Total
70 - 75	4766.0	3.0	0.0	4769.0
80 - 85	7341.0	0.0	0.0	7341.0
90 - 95	3541.0	0.0	0.0	3541.0

40 - 45	0.0	1095.0	0.0	1095.0
50 - 55	0.0	1685.0	0.0	1685.0
60 - 65	0.0	3322.0	0.0	3322.0
0 - 5	0.0	0.0	37.0	37.0
10 - 15	0.0	0.0	135.0	135.0
20 - 25	0.0	0.0	259.0	259.0
30 - 35	0.0	0.0	711.0	711.0
Total	15648.0	6105.0	1142.0	22895.0

```
[28]: oci_plt = oci_cond['Total'][0:8].sort_values(ascending=False)
oci_plt.plot(kind='bar', width=0.90)
plt.title('Street Conditions by Index Range')
plt.xlabel('Index Range')
plt.ylabel('# of Streets')
plt.show()
```



Summary Statistics

```
[29]: # summary statistics
summ_stats = pd.DataFrame(df['oci'].describe()).T
summ_stats
```

```
[29]:
```

	count	mean	std	min	25%	50%	75%	max
oci	23005.0	74.791413	16.784048	0.0	66.3	79.06	87.3	100.0

```
[30]: IQR = summ_stats['75%'][0] - summ_stats['25%'][0]
low_outlier = summ_stats['25%'][0] - 1.5*(IQR)
high_outlier = summ_stats['75%'][0] + 1.5*(IQR)

print('Low Outlier:', low_outlier)
print('High Outlier:', high_outlier)
```

Low Outlier: 34.8
High Outlier: 118.8

```
[31]: print("\033[1m" + 'Overall Condition Index (OCI) Summary' + "\033[1m")
def oci_by_range():
    pd.options.display.float_format = '{:,.2f}'.format
    new = df.groupby('OCI Range')['oci']\
        .agg(["mean",
              "median",
              "std",
              "min",
              "max"])

    new.loc['Total'] = new.sum(numeric_only=True, axis=0)
    column_rename = {'mean': 'Mean', 'median': 'Median',
                     'std': 'Standard Deviation',
                     'min': 'Minimum', 'max': 'Maximum'}

    dfsummary = new.rename(columns = column_rename)
    return dfsummary

oci_by_range = oci_by_range()
oci_by_range
```

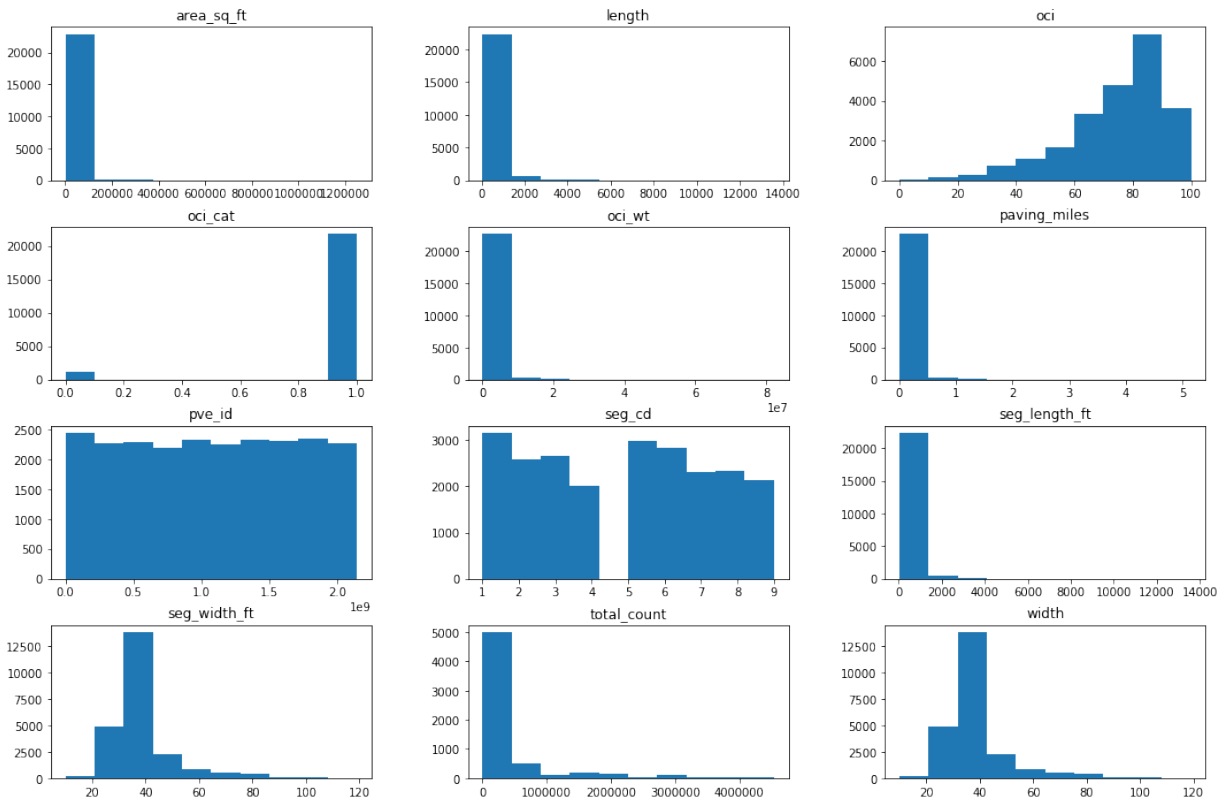
Overall Condition Index (OCI) Summary

```
[31]:
```

	Mean	Median	Standard Deviation	Minimum	Maximum
OCI Range					
0 - 5	6.13	8.00	3.70	0.00	9.69
10 - 15	15.66	16.40	2.82	10.11	19.84
20 - 25	25.77	26.17	2.91	20.12	29.96
30 - 35	35.63	36.04	2.80	30.04	39.98
40 - 45	45.37	45.58	2.88	40.00	49.98
50 - 55	55.62	56.00	2.88	50.00	59.98
60 - 65	65.56	65.80	2.82	60.00	69.99
70 - 75	75.11	75.16	2.97	70.00	79.99
80 - 85	85.14	85.15	2.84	80.00	89.99
90 - 95	93.44	92.89	2.57	90.00	99.33
Total	503.42	507.19	29.18	450.27	548.73

Histogram Distributions

```
[32]: # histograms
df.hist(grid=False, figsize=(18,12))
plt.show()
```



Boxplot Distribution (OCI)

```
[33]: # selected boxplot distribution for oci values
print("\033[1m" + 'Boxplot Distribution' + "\033[1m")

# Boxplot of age as one way of showing distribution
fig = plt.figure(figsize = (10,1.5))
plt.title ('Boxplot: Overall Condition Index (OCI)')
plt.xlabel('Speed Limit')
plt.ylabel('Value')


sns.boxplot(data=df['oci'],
            palette="coolwarm",
            orient='h',
            linewidth=2.5)

plt.show()

IQR = summ_stats['75%'][0] - summ_stats['25%'][0]

print('The first quartile is %s. '%summ_stats['25%'][0])
print('The third quartile is %s. '%summ_stats['75%'][0])
print('The IQR is %s. '%round(IQR,2))
print('The mean is %s. '%round(summ_stats['mean'][0],2))
print('The standard deviation is %s. '%round(summ_stats['std'][0],2))
print('The median is %s. '%round(summ_stats['50%'][0],2))
```

Boxplot Distribution



The first quartile is 66.3.

The third quartile is 87.3.

The IQR is 21.0.

The mean is 74.79.

The standard deviation is 16.78.

The median is 79.06.

Correlation Matrix

```
[34]: # assign correlation function to new variable

corr = df.corr()

# for triangular matrix


matrix = np.triu(corr)

plt.figure(figsize=(
    10,10
))

# parse corr variable into triangular matrix

sns.heatmap(df.corr(
    method='pearson'),
    annot=True,
    linewidths=.5,
    cmap="coolwarm",
    mask=matrix,
    square = True,
    cbar_kws={'label': 'Correlation Index'},
    vmin=-1,
    vmax=1
))

plt.show()
```



Multicollinearity

Let us narrow our focus by removing highly correlated predictors and passing the rest into a new dataframe.

```
[35]: cor_matrix = df.corr().abs()

upper_tri = cor_matrix.where(np.triu(np.ones(cor_matrix.shape),
                                         k=1).astype(np.bool))

to_drop = [column for column in upper_tri.columns if
            any(upper_tri[column] > 0.75)]

print('These are the columns prescribed to be dropped: %s'%to_drop)
```

These are the columns prescribed to be dropped: ['area_sq_ft', 'oci_wt', 'length', 'width', 'paving_miles']

Pre-Processing

Based on the prescribed output of the multicollinearity outcome, we should remove `area_sq_ft`, `oci_wt`, `length`, `width`, `paving_miles`, respectively. However, area in square feet is derived from length (x) width values, and converted to paving miles. Removing all of these features is not necessary. We can keep area in square feet, as long as we remove the rest.

Feature Engineering

The start date is subtracted from the end date and converted to number of days as one column.

```
[36]: df['date_end'] = pd.to_datetime(df['date_end'])
      df['date_start'] = pd.to_datetime(df['date_start'])

      # 7 rows with missing values are dropped in the following line
      day_diff = df.dropna(subset=['date_end',
                                   'date_start'],
                           inplace=True)

      df['day_diff'] = (df['date_end'] - df['date_start']).dt.days.astype(int)
```

```
[37]: zero_days = df['day_diff'].value_counts()[0]
      percent_days = round(zero_days/len(df), 2)*100
      print('There are', zero_days, 'rows with "0".')
      print('That is roughly', percent_days, '% of the data.')
```

There are 18451 rows with "0".

That is roughly 80.0 % of the data.

The residential, collector, major, prime, local, and alley functional classes are converted to dummy variables.

```
[38]: df['func_class'].value_counts()
      df['func_cat'] = df['func_class'].map({'Residential': 1,
                                             'Collector': 2,
                                             'Major': 3, 'Prime':4,
                                             'Local':5, 'Alley':6})
```

The AC Improved, PCC Jointed Concrete, AC Unimproved, and UnSurfaced pavement classes are converted to dummy variables.

```
[39]: df['pvm_class'].value_counts()
      df['pvm_cat'] = df['pvm_class'].map({'AC Improved': 1,
                                             'PCC Jointed Concrete': 2,
                                             'AC Unimproved': 3,
                                             'UnSurfaced':4})
```

The current status of the job (i.e., post construction, design, bid/award, construction, and planning) is also converted to dummy variables.

```
[40]: df['status'].value_counts()
      df['status_cat'] = df['status'].map({'post construction': 1,
                                             'design': 2,
                                             'bid / award': 3,
                                             'construction':4,
                                             'planning': 5})
```

Dropping Non-Useful/Re-classed Columns

Columns with explicit titles (i.e., names) and non-convertible/non-meaningful strings are dropped. Redundant columns (columns that have been cast to dummy variables) have also been dropped in conjunction with the index column which serves no purpose for this experiment.

```
[41]: # drop unnecessary columns
df = df.drop(columns=['street_from',
                      'street_to',
                      'street_name',
                      'seg_id',
                      'street',
                      'pve_id',
                      'title',
                      'project_manager',
                      'project_manager_phone',
                      'project_id',
                      'resident_engineer',
                      'address_street',
                      'date_moratorium',
                      'OCI Range',
                      'total_count'])

df = df.reset_index(drop=True)

# drop variables exhibiting multicollinearity
df = df.drop(columns=['seg_length_ft',
                      'seg_width_ft',
                      'length',
                      'width',
                      'paving_miles',
                      'oci_wt'])

# drop re-classed columns
df = df.drop(columns=['func_class',
                      'pvm_class',
                      'status',
                      'type',
                      'date_end',
                      'date_start',
                      'oci_desc'])
```

The original dataframe is copied into a new dataframe *df1* in order to continue the final steps in the pre-processing endeavor. This is to avoid any mis-steps or adverse/unintended effects on the original dataframe.

```
[42]: # create new dataframe for final pre-processing steps
df1 = df.copy()
```

One consequence of pre-processing data is that additional missing values may be brought into the mix, so one final sanity check for this phenomenon is commenced as follows.

```
[43]: df_check = df.isna().sum()
df_check[df_check>0]
```

```
[43]: Series([], dtype: int64)
```

```
[44]: cor_matrix = df.corr().abs()
upper_tri = cor_matrix.where(np.triu(np.ones(cor_matrix.shape),
                                         k=1).astype(np.bool))

to_drop = [column for column in upper_tri.columns if
            any(upper_tri[column] > 0.75)]

print('These are the columns prescribed to be dropped: %s'%to_drop)
```

These are the columns prescribed to be dropped: []

Handling Class Imbalance

Multiple methods for balancing a dataset exist like “undersampling the majority classes” (Fregly & Barth, 2021, p. 178). To account for the large gap (95%) of mis-classed data on the “poor” condition class, “oversampling the minority class up to the majority class” (p. 179) is commenced. However, such endeavor cannot proceed in good faith without the unsupervised dimensionality reduction technique of Principal Component Analysis (PCA), which is carried out “to compact the dataset and eliminate irrelevant Features” (Naseriparsa & Kashani, 2014, p. 33). In this case, a new dataframe is reduced down into the first two principal components since the largest percent variance explained exists therein.

```
[45]: # the first two principal components are used
pca = PCA(n_components=2, random_state=777)
data_2d = pd.DataFrame(pca.fit_transform(df1.iloc[:,0:9]))
```

The dataframe is prepared for scatterplot analysis as follows.


```
[46]: data_2d = pd.concat([data_2d, df1['oci_cat']], axis=1)
data_2d.columns = ['x', 'y', 'oci_cat']; data_2d
```

```
[46]:
```

	x	y	oci_cat
0	7,986.11	-38.95	0
1	47.96	-40.17	1
2	-435.67	-39.63	1
3	221.51	-39.71	1
4	-348.92	-39.08	0
...
22993	-15,801.67	-40.53	1
22994	12,768.33	114.26	1
22995	9,128.33	114.06	1
22996	-12,991.67	-40.48	1
22997	-12,800.19	-40.25	1

[22998 rows x 3 columns]

```
[47]: sns.lmplot('x', 'y', data_2d,
                fit_reg=False,
                hue='oci_cat',
                palette=['#00BFC4',
                        '#F8766D'])
plt.title('Class Imbalance in Street Condition Index'); plt.show()
```



The dataset is oversampled into a new dataframe *df2*.

The adaptive synthetic sampling approach (ADAYSN) is leveraged “where more synthetic data is generated for minority class examples that are harder to learn compared to those minority examples that are easier to learn” (He et al., 2008). This allows for the minority class to be more closely matched (up-sampled) to the majority class for an approximately even 50/50 weight distribution.

```
[48]: ada = ADASYN(random_state=777)
      X_resampled, y_resampled = ada.fit_resample(df1.iloc[:,0:7],
                                                df1['oci_cat'])
```

```
[49]: df2 = pd.concat([pd.DataFrame(X_resampled),
                      pd.DataFrame(y_resampled)], axis=1)
      df2.columns = df1.columns
```

The classes are re-balanced in a new dataframe using oversampling:

```
[50]: # rebalanced classes in new df
      df2['oci_cat'].value_counts()
      zero_count = df2['oci_cat'].value_counts()[0]
      one_count = df2['oci_cat'].value_counts()[1]
      zero_plus_one = zero_count + one_count
```

```

print('Poor Condition Size:', zero_count)
print('Good Condition Size:', one_count)
print('Total Condition Size:', zero_plus_one)
print('Percent in Poor Condition:', round(zero_count/zero_plus_one,2))
print('Percent in Good Condition:', round(one_count/zero_plus_one,2))

```

```

Poor Condition Size: 21714
Good Condition Size: 21858
Total Condition Size: 43572
Percent in Poor Condition: 0.5
Percent in Good Condition: 0.5

```

The dataframe can now be prepared as a flat .csv file if so desired.

Train-Test-Validation Split

```

[51]: #Divide train set by .7, test set by .15, and valid set .15
size_train = 30500
size_valid = 6536
size_test = 6536
size_total = size_test + size_valid + size_train
train, test = train_test_split(df2, train_size = size_train,\
                               random_state = 777)
valid, test = train_test_split(test, train_size = size_valid,\
                               random_state = 777)

print('Training size:', size_train)
print('Validation size:', size_valid)
print('Test size:', size_test)
print('Total size:', size_train + size_valid + size_test)
print('Training percentage:', round(size_train/(size_total),2))
print('Validation percentage:', round(size_valid/(size_total),2))
print('Test percentage:', round(size_test/(size_total),2))

```

```

Training size: 30500
Validation size: 6536
Test size: 6536
Total size: 43572
Training percentage: 0.7
Validation percentage: 0.15
Test percentage: 0.15

```

```

[52]: # define (list) the features
X_var = list(df2.columns)

# define the target
target = 'oci_cat'
X_var.remove(target)
X_train = train[X_var]
y_train = train[target]
X_test = test[X_var]
y_test = test[target]

```

```
X_valid = valid[X_var]
y_valid = valid[target]
```

```
[53]: # rearrange columns so that the target column is set up first
# for later training
df2 = df2[['oci_cat', 'oci', 'area_sq_ft', 'seg_cd', 'day_diff',
           'func_cat', 'pvm_cat', 'status_cat']]
```

```
[54]: # reinspect the dataframe
df2.head()
```

```
[54]:
```

	oci_cat	oci	area_sq_ft	seg_cd	day_diff	func_cat	pvm_cat	status_cat
0	0	34.14	28,137.78	2	0	6	1	0
1	1	97.25	20,199.63	2	0	6	2	1
2	1	62.67	19,716.00	2	0	6	2	1
3	1	68.86	20,373.18	9	0	6	2	1
4	0	28.67	19,802.75	9	0	6	2	0

Transfer The Final Dataframe (df2) to S3 Bucket

```
[55]: s3_client = boto3.client("s3")
BUCKET='waterteam1'
KEY='raw_files/df2/df2.csv'
response = s3_client.get_object(Bucket=BUCKET, Key=KEY)

with io.StringIO() as csv_buffer:
    df2.to_csv(csv_buffer, index=False, header=True)

    response = s3_client.put_object(
        Bucket=BUCKET, Key=KEY, Body=csv_buffer.getvalue()
    )
```

Modeling and Training

Logistic Regression

Herein, the classical Anaconda-based scikit-learn approach is leveraged to train the logistic regression model on the validation set.

```
[56]: # Un-Tuned Logistic Regression Model
logit_reg = LogisticRegression(random_state=777)
logit_reg.fit(X_train, y_train)

# Predict on validation set
logit_reg_pred1 = logit_reg.predict(X_valid)

# accuracy and classification report (Untuned Model)
print('Untuned Logistic Regression Model')
print('Accuracy Score')
print(accuracy_score(y_valid, logit_reg_pred1))
print('Classification Report \n',
      classification_report(y_valid, logit_reg_pred1))
```

Untuned Logistic Regression Model

Accuracy Score

0.8959608323133414

Classification Report

	precision	recall	f1-score	support
0	0.93	0.85	0.89	3286
1	0.86	0.94	0.90	3250
accuracy			0.90	6536
macro avg	0.90	0.90	0.90	6536
weighted avg	0.90	0.90	0.90	6536

Next, the logistic regression model is tuned using `RandomizedSearchCV()` and cross validated using repeated stratified kfold with five splits and two repeats. A set of hyperparameters are subsequently defined to produce an overall best accuracy score in conjunction with a set of optimal hyperparameters.

```
[57]: model1 = LogisticRegression(random_state=777)
cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=2,
                             random_state=777)

space = dict()

# define search space
space['solver'] = ['newton-cg', 'lbfgs', 'liblinear']
space['penalty'] = ['none', 'l1', 'l2', 'elasticnet']
space['C'] = loguniform(1e-5, 100)

# define search
search = RandomizedSearchCV(model1, space,
                             scoring='accuracy',
                             n_jobs=-1, cv=cv, random_state=777)

# execute search
result = search.fit(X_train, y_train)

# summarize result
print('Best Score: %s' % result.best_score_)
print('Best Hyperparameters: %s' % result.best_params_)
```

Best Score: 0.9518524590163935

Best Hyperparameters: {'C': 0.005639439254142048, 'penalty': 'l2', 'solver': 'lbfgs'}

Training, Testing, and Deploying a Model with Amazon SageMaker's Built-in XGBoost Model

```
[58]: # Define IAM role
role = get_execution_role()

# set the region of the instance
my_region = boto3.session.Session().region_name

# this line automatically looks for the XGBoost image URI and
```

```
# builds an XGBoost container.
xgboost_container = sagemaker.image_uris.retrieve("xgboost",
                                                  my_region,
                                                  "latest")

print("Success - the MySageMakerInstance is in the " + my_region + \
      " region. You will use the " + xgboost_container + \
      " container for your SageMaker endpoint.")
```

Success - the MySageMakerInstance is in the us-east-1 region. You will use the 811284229777.dkr.ecr.us-east-1.amazonaws.com/xgboost:latest container for your SageMaker endpoint.

```
[59]: train, test = np.split(df2.sample(frac=1, random_state=777),
                               [int(0.7 * len(df2))])

print(train.shape, test.shape)
```

(30500, 8) (13072, 8)

Transfer The Training Data to S3 Bucket

```
[60]: s3_client = boto3.client("s3")

BUCKET='waterteam1'
KEY='raw_files/train/train.csv'

response = s3_client.get_object(Bucket=BUCKET, Key=KEY)

with io.StringIO() as csv_buffer:
    train.to_csv(csv_buffer, index=False, header=False)

    response = s3_client.put_object(
        Bucket=BUCKET, Key=KEY, Body=csv_buffer.getvalue()
    )
```

```
[61]: # input training parameters
s3_input_train = sagemaker.inputs.TrainingInput(s3_data=\
        's3://{}/raw_files/train'.format(BUCKET), content_type='csv')
```

Setting up the SageMaker Session and Supplying Instance for XGBoost Model

```
[62]: sess = sagemaker.Session()
xgb = sagemaker.estimator.Estimator(xgboost_container,role,
                                    instance_count=1,
                                    instance_type='ml.m5.large',
                                    output_path='s3://{}/output'.format(BUCKET),
                                    sagemaker_session=sess)

# parse in the hyperparameters
xgb.set_hyperparameters(max_depth=5,eta=0.2,gamma=4,min_child_weight=6,
                        subsample=0.8,silent=0,
                        objective='binary:logistic',num_round=100)
```


Train The Model

```
[63]: xgb.fit({'train': s3_input_train})
```

```
2022-04-10 22:17:12 Starting - Starting the training job...
2022-04-10 22:17:29 Starting - Preparing the instances for
trainingProfilerReport-1649629032: InProgress
...
2022-04-10 22:18:55 Downloading - Downloading input data...
2022-04-10 22:19:56 Training - Downloading the training image..Arguments: train
[2022-04-10:22:20:25:INFO] Running standalone xgboost training.
[2022-04-10:22:20:25:INFO] Path /opt/ml/input/data/validation does not exist!
[2022-04-10:22:20:25:INFO] File size need to be processed in the node: 1.09mb.

Available memory size in the node: 294.12mb
[2022-04-10:22:20:25:INFO] Determined delimiter of CSV input is ','
[22:20:25] S3DistributionType set as FullyReplicated
[22:20:25] 30500x7 matrix with 213500 entries loaded from

/opt/ml/input/data/train?format=csv&label_column=0&delimiter=,
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[0]#011train-error:3.3e-05
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[1]#011train-error:3.3e-05
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[2]#011train-error:0
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[3]#011train-error:0
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[4]#011train-error:0
[22:20:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[5]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[6]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[7]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[8]#011train-error:0
```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[9]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[10]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[11]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[12]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[13]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[14]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[15]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[16]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[17]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[18]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[19]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[20]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[21]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[22]#011train-error:0

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[23]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[24]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[25]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[26]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[27]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[28]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[29]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[30]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[31]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[32]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[33]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[34]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 2 extra nodes, 0
pruned nodes, max_depth=1
[35]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[36]#011train-error:0

```

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[37]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[38]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[39]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[40]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[41]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[42]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[43]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[44]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[45]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[46]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[47]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[48]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[49]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[50]#011train-error:0

```

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[51]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[52]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[53]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[54]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[55]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[56]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[57]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[58]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[59]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[60]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[61]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[62]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[63]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[64]#011train-error:0

```

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[65]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[66]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[67]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[68]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[69]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[70]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[71]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[72]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[73]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[74]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[75]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[76]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[77]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[78]#011train-error:0

```

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[79]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[80]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[81]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[82]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[83]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[84]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[85]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[86]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[87]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[88]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[89]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[90]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[91]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[92]#011train-error:0

```

```

[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[93]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[94]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[95]#011train-error:0
[22:20:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[96]#011train-error:0
[22:20:27] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[97]#011train-error:0
[22:20:27] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[98]#011train-error:0
[22:20:27] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 0 extra nodes, 0
pruned nodes, max_depth=0
[99]#011train-error:0

```

```

2022-04-10 22:20:43 Uploading - Uploading generated training model
2022-04-10 22:20:43 Completed - Training job completed
Training seconds: 118
Billable seconds: 118

```

Deploying The Predictor

```

[64]: xgb_predictor = xgb.deploy(initial_instance_count=1,
                                instance_type='ml.m5.large')

```

-----!

Running Predictions

```

[65]: from sagemaker.serializers import CSVSerializer

# load the data into an array
test_array = test.drop(['oci_cat'], axis=1).values

# set the serializer type
xgb_predictor.serializer = CSVSerializer()

# predict!
predictions = xgb_predictor.predict(test_array).decode('utf-8')

# and turn the prediction into an array

```



```
predictions_array = np.fromstring(predictions[1:], sep=',')
print(predictions_array.shape)
```

(13072,)

Evaluating The Model

```
[66]: cm = pd.crosstab(index=test['oci_cat'],
                        columns=np.round(predictions_array),
                        rownames=['Observed'],
                        colnames=['Predicted'])
tn = cm.iloc[0,0]; fn = cm.iloc[1,0]; tp = cm.iloc[1,1];
fp = cm.iloc[0,1]; p = (tp+tn)/(tp+tn+fp+fn)*100
print("\n{0:<20}{1:<4.1f}%\n".format("Overall Classification Rate: ", p))
print("{0:<15}{1:<15}{2:>8}".format("Predicted", "Poor Condition",
                                   "Good Condition"))

print("Observed")
print("{0:<15}{1:<2.0f}% ({2:<}){3:>6.0f}% ({4:<})".format("Poor Condition", \
                                                         tn/(tn+fn)*100,tn, fp/(tp+fp)*100, fp))
print("{0:<16}{1:<1.0f}% ({2:<}){3:>7.0f}% ({4:<}) \n".format("Good Condition", \
                                                            fn/(tn+fn)*100,fn, tp/(tp+fp)*100, tp))
```

Overall Classification Rate: 100.0%

Predicted	Poor Condition	Good Condition
Observed		
Poor Condition	100% (6497)	0% (0)
Good Condition	0% (0)	100% (6575)

Terminating the Endpoint To Save on Costs

```
[67]: # clean-up by deleteting endpoint
xgb_predictor.delete_endpoint(delete_endpoint_config=True)
```

References

Amazon Web Services. (n.d.). *Amazon Athena*.

<https://aws.amazon.com/athena/?whats-new-cards.sort-by=item.additionalFields.postDateTime&whats-new-cards.sort-order=desc>

Amazon Web Services. (n.d.). *Build, train, and deploy a machine learning model with Amazon SageMaker*.

<https://aws.amazon.com/getting-started/hands-on/build-train-deploy-machine-learning-model-sagemaker/>

Fregly, C., & Barth, A. (2021). *Data Science on AWS*. O'Reilly.

Garrick, D. (2021, September 12). San Diego to spend \$700K assessing street conditions to spend repair money wisely. *The San Diego Union-Tribune*.

<https://www.sandiegouniontribune.com/news/politics/story/2021-09-12/san-diego-to-spend-700k-assessing-street-conditions-to-spend-repair-money-wisely>

He, H., Bai, Y., Garcia, E. & Li, S. (2008). ADASYN: Adaptive synthetic sampling approach for imbalanced learning.

2008 IEEE International Joint Conference on Neural Networks (IEEE World Congress on Computational

Intelligence),1322-1328.

<https://ieeexplore.ieee.org/document/4633969>

Naseriparsa, M. & Kashani, M.M.R. (2014). Combination of PCA with SMOTE Resampling to Boost the Prediction Rate in Lung Cancer Dataset.

International Journal of Computer Applications, 77(3) 33-38. <https://doi.org/10.5120/13376-0987>